

Tutorial Introduction to Neural Networks with an eye towards linguistic applications

Shane Steinert-Threlkeld

January 25, 2019



Today's Plan

1. Neural Networks: computation
2. Neural Networks: learning
3. Hands-on experiment: learning quantifiers
4. Some practical tips
5. Further Topics + Resources

What I'm Presupposing

Some mathematical notation/concepts from:

- Linear algebra (matrix multiplication, e.g.)
- Multivariate calculus (partial derivatives)

What I'm Presupposing

Some mathematical notation/concepts from:

- Linear algebra (matrix multiplication, e.g.)
- Multivariate calculus (partial derivatives)

Some programming experience:

- Basics of Python
- Basic syntax in NumPy

What I'm Presupposing

Some mathematical notation/concepts from:

- Linear algebra (matrix multiplication, e.g.)
- Multivariate calculus (partial derivatives)

Some programming experience:


- Basics of Python
- Basic syntax in NumPy

But: no formal requirements; all concepts and syntax can be explained intuitively, so please ask for clarification at all points!


Materials: Slides + Jupyter Notebook

<https://github.com/shanest/nn-tutorial>


Artificial Neuron




Artificial Neuron



Artificial Neuron




Artificial Neuron




$$a = f(a_0 \cdot w_0 + a_1 \cdot w_1 + a_2 \cdot w_2)$$

Artificial Neuron



$$a = f(a_0 \cdot w_0 + a_1 \cdot w_1 + a_2 \cdot w_2)$$

Activation Function




More on choosing activation functions later in the tutorial.


Computing ‘and’

p	q	$p \wedge q$
1	1	1
1	0	0
0	1	0
0	0	0

Computing 'and'




Computing 'and'



$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$


Computing 'and'



$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

Computing 'and'




$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

$$a = \sigma(0 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

Computing 'and'




$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$


$$a = \sigma(0 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

$$a = \sigma(0 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-30) \approx 0$$

Computing 'and'




Computing 'and'




Computing 'xor'

p	q	$p \text{ xor } q$
1	1	0
1	0	1
0	1	1
0	0	0


Computing 'xor'




Computing 'xor'




Computing 'xor'




Computing 'xor'



Computing 'xor'




Computing 'xor'




xor is not *linearly separable*


Computing 'xor'




Computing 'xor'




Computing 'xor'




Computing 'xor'




Computing 'xor'



Computing 'xor'




Computing 'xor'




Exercise: show that the hidden units behave as labeled.


Computing 'xor'



Computing 'xor'



Computing 'xor'



Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\begin{aligned}\vec{a} &= f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right) \\ &= f(xW + b)\end{aligned}$$

Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\begin{aligned}\vec{a} &= f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right) \\ &= f(xW + b)\end{aligned}$$

- x_j^i : j th feature of input i
- w_k^l : weight from neuron k to neuron l in next layer
- b_m : bias to neuron m in next layer

Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\vec{a} = f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right)$$

$$= f(xW + b)$$

- x_j^i : j th feature of input i
- w_k^l : weight from neuron k to neuron l in next layer
- b_m : bias to neuron m in next layer

Exercises:

- write down W^1 and W^2 for the xor network.
- re-write the above as $f(xW)$ by adding a column of 1s to x and a new row to W .

Hidden Representations

Key idea: hidden layers of a neural network can encode high-level/abstract features of the input.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.
- Update the weights and biases to move the network's outputs closer to the true outputs.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.
- Update the weights and biases to move the network's outputs closer to the true outputs.

The last step is done via *gradient descent* (and refinements thereof).

Gradient Descent: Example

Task: predict a true value $y = 2$.

Gradient Descent: Example

Task: predict a true value $y = 2$.

“Model”: one parameter θ , outputs $\hat{y} = \theta$.

Gradient Descent: Example


Task: predict a true value $y = 2$.

“Model”: one parameter θ , outputs $\hat{y} = \theta$.

Loss function:

$$\mathcal{L}(\theta, y) = (\hat{y}(\theta) - y)^2$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$

Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:

$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2) \cdots)W^L + b^L)$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:

$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2)\cdots)W^L + b^L)$$

All of the weights and biases form a long vector of parameters θ . So instead of a partial derivative, we take a *gradient*:

$$\nabla_{\theta} \mathcal{L}(\hat{y}(\theta), y) = \left\langle \frac{\partial}{\partial \theta_1} \mathcal{L}, \dots, \frac{\partial}{\partial \theta_N} \mathcal{L} \right\rangle$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:


$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2) \cdots)W^L + b^L)$$

All of the weights and biases form a long vector of parameters θ . So instead of a partial derivative, we take a *gradient*:


$$\nabla_{\theta} \mathcal{L}(\hat{y}(\theta), y) = \left\langle \frac{\partial}{\partial \theta_1} \mathcal{L}, \dots, \frac{\partial}{\partial \theta_N} \mathcal{L} \right\rangle$$

The (negative) gradient tells us *which direction in 'parameter space'* to walk in order to make the loss (\mathcal{L}) smaller, i.e. to make the network's output closer to the true output.


Activation Functions



Activation Functions



Activation Functions



ReLU is incredibly popular at the moment, as are refinements: leaky ReLU, exponential linear (ELU), gaussian linear (GLU), ...

Loss Functions

Regression: different kinds of geometrical distances, e.g.

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Loss Functions

Regression: different kinds of geometrical distances, e.g.

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Classification: *cross entropy*:

$$\ell(\hat{y}, y) = - \sum y_i \cdot \ln(\hat{y}_i) = - \ln(\hat{y}_{\text{true class}})$$


Loss Functions

Regression: different kinds of geometrical distances, e.g.

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$

Classification: *cross entropy*:

$$\ell(\hat{y}, y) = - \sum y_i \cdot \ln(\hat{y}_i) = - \ln(\hat{y}_{\text{true class}})$$



Learned Representations

Key idea: a neural network can learn *which* high-level/abstract features of the input are useful in helping it solve its task. (Features are learned, instead of engineered by us.)

Anatomy of a DL Experiment

① Specify parameters

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!
 - ① Evaluate at regular intervals
 - ② Measure your variables of interest
 - ③ Monitor train/test loss
 - ④ Early stopping [later in tutorial]
- ⑤ Analyze
 - quantitative
 - qualitative behavior

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!
 - ① Evaluate at regular intervals
 - ② Measure your variables of interest
 - ③ Monitor train/test loss
 - ④ Early stopping [later in tutorial]
- ⑤ Analyze
 - quantitative
 - qualitative behavior

NOTE: keep detailed records about what you're doing!

For guidance on keeping records and writing up results, see [these lecture notes](#) from Sam Bowman.

To the code!

<https://github.com/shanest/nn-tutorial/blob/master/tutorial.ipynb>

Balancing the Data

For classification tasks, it's very important to *balance the training data*, so that it contains (roughly) equal numbers of examples for each class. Otherwise, a network can very quickly get stuck in a local minimum, where it uniformly guesses the most-frequent class.

Two methods:

- ① Downsampling: randomly sample as many examples from each class as in your least-frequent class
- ② Upsampling: repeat examples from your less-frequent classes until you have as many as the most-frequent

Early Stopping

How do you know how many epochs to train for? One common method:
A LOT, but put in a condition for *early stopping*.

Over-fitting: detection

Neural networks, especially very large, deep ones, run the risk of *over-fitting* their training data. (Sometimes, this is referred to as “memorizing” the training data.)

Over-fitting: avoidance (regularization)

Two very popular and successful *regularization* techniques to combat over-fitting (and generally make your life better):

- Dropout (Srivastava et al. 2014): randomly ‘turn off’ (set to 0) a certain percentage of input nodes to this layer
- Batch normalization (Ioffe and Szegedy 2015): scale the inputs to the layer so that they’re roughly averaged around 0 and normally distributed.

(I’ve found BN *very* useful in my own research.)

Note: `.eval()` and `.train()` on a PyTorch `nn.Module` turn these sorts of methods off/on based on whether you are in training or ‘inference’ / prediction mode.

Hyper-parameter Tuning


How do you decide how to set the parameters of your experiment? There are so many knobs to turn! (Listed roughly in order of importance of turning them.)

- ① Network architecture: depth (number of layers) and width (size of layers)
[or even different network types altogether]
- ② Activation functions
- ③ Optimizers
 - Learning rates + other parameters here
- ④ ...

Hyper-parameter Tuning (cont.)

Main idea: try a whole bunch of settings! Choose the setting that performs best. Some notes:

- This requires a third set, a *development (dev)* set, in addition to training/testing.
You choose the best hyper-parameters based on best performance (however measured) on the *dev* set. Then you evaluate that model on the *test* set.
- *Random* search in parameter space better than 'grid' search:



Bergstra and Bengio 2012

Scaling Up

Once models and datasets become non-trivially sized, your personal computer likely won't suffice to run experiments.

A common paradigm: prototype locally, experiment externally on a 'cluster'.

Two most critical components:

- Large RAM for storing big models and datasets
- GPUs (Graphics Processing Units)!
These are made to massively parallelize linear algebra of the kind computed by networks, so can dramatically speed up training and inference.

If you can't access these things via Cornell, you can buy compute time on appropriate machines via Amazon Web Services or Google Cloud Platform.

References I



Bergstra, James and Yoshua Bengio (2012). “Random Search for Hyper-Parameter Optimization”. In: *Journal of Machine Learning Research* 13, pp. 281–305.



Ioffe, Sergey and Christian Szegedy (2015). “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”. In: arXiv: 1502.03167. URL: <http://arxiv.org/abs/1502.03167>.



Srivastava, Nitish et al. (2014). “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”. In: *Journal of Machine Learning Research* 15, pp. 1929–1958. DOI: 10.1214/12-AOS1000.