

Tutorial Introduction to Neural Networks with an eye towards linguistic applications

Shane Steinert-Threlkeld

January 25, 2019




Today's Plan


Materials: Slides + Jupyter Notebook

<https://github.com/shanest/nn-tutorial>


Artificial Neuron




Artificial Neuron



Artificial Neuron




Artificial Neuron



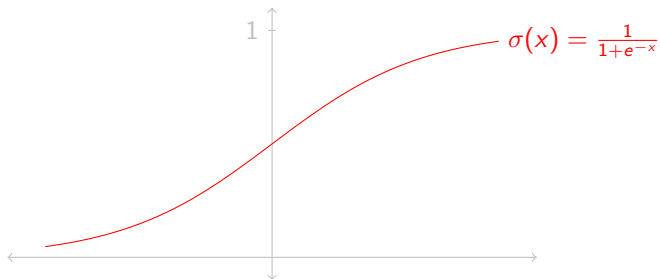
$$a = f(a_0 \cdot w_0 + a_1 \cdot w_1 + a_2 \cdot w_2)$$

Artificial Neuron



$$a = f(a_0 \cdot w_0 + a_1 \cdot w_1 + a_2 \cdot w_2)$$

Activation Function




More on choosing activation functions later in the tutorial.

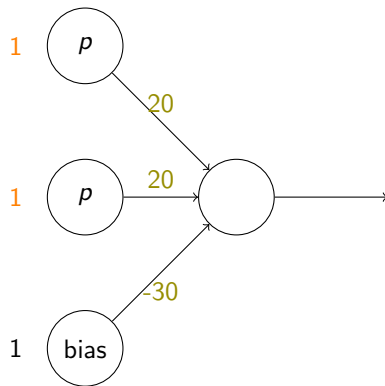
Computing ‘and’

p	q	$p \wedge q$
1	1	1
1	0	0
0	1	0
0	0	0

Computing 'and'




Computing 'and'



$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

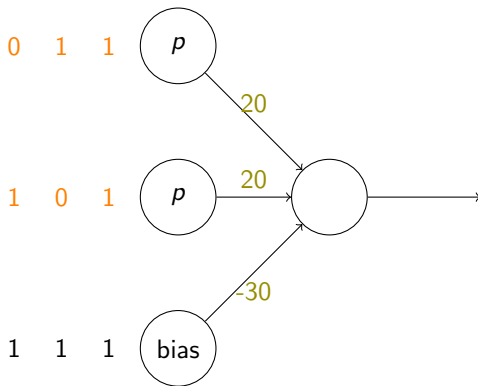
Computing 'and'



$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

Computing 'and'

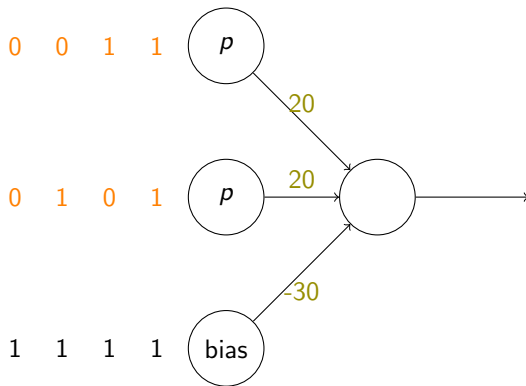


$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

$$a = \sigma(0 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

Computing 'and'




$$a = \sigma(1 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(10) \approx 1$$

$$a = \sigma(1 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

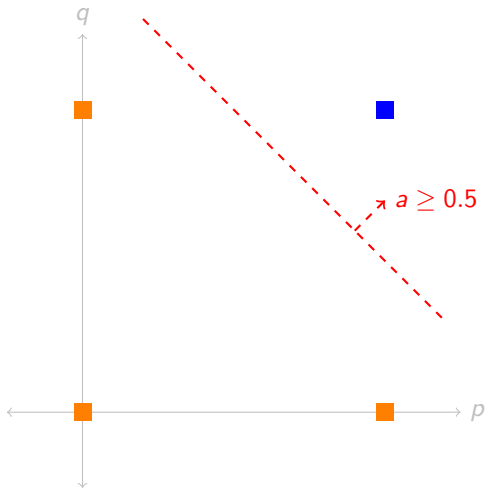
$$a = \sigma(0 \cdot 20 + 1 \cdot 20 + 1 \cdot -30) = \sigma(-10) \approx 0$$

$$a = \sigma(0 \cdot 20 + 0 \cdot 20 + 1 \cdot -30) = \sigma(-30) \approx 0$$

Computing 'and'



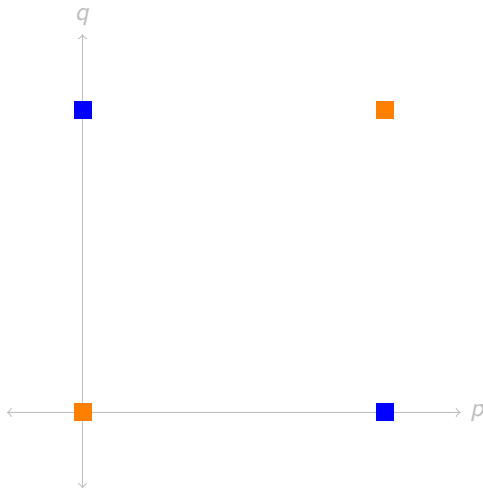
Computing 'and'



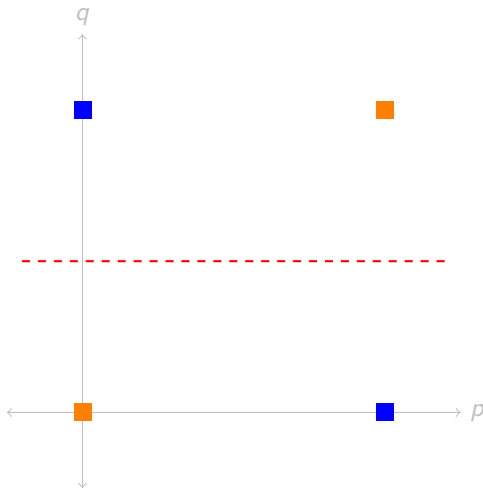
Computing 'xor'

p	q	$p \text{ xor } q$
1	1	0
1	0	1
0	1	1
0	0	0

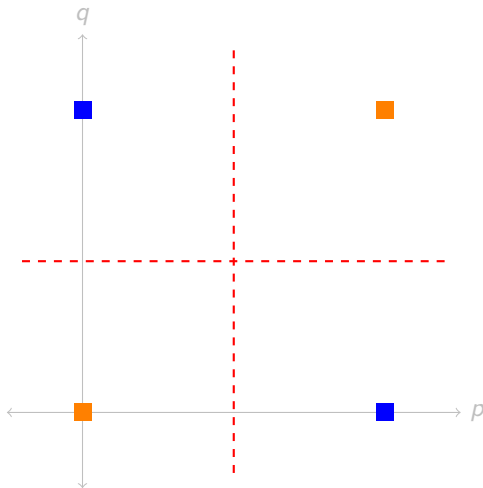
Computing 'xor'



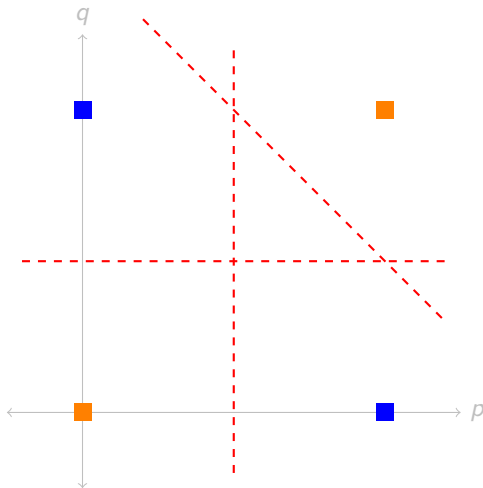
Computing 'xor'



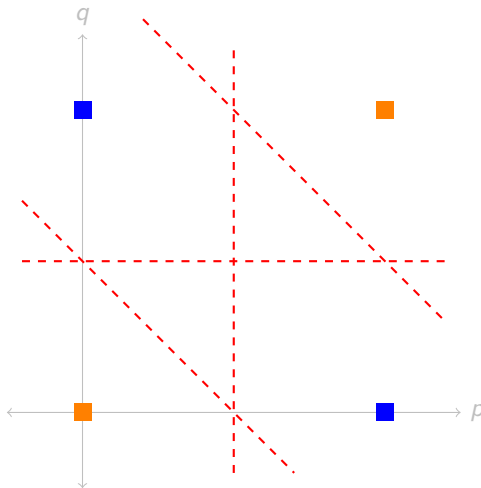
Computing 'xor'




Computing 'xor'



Computing 'xor'




Computing 'xor'

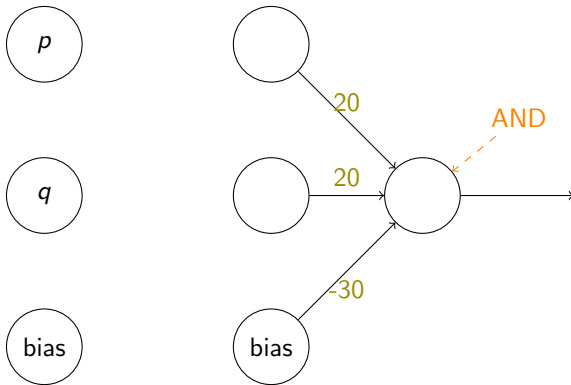


xor is not *linearly separable*

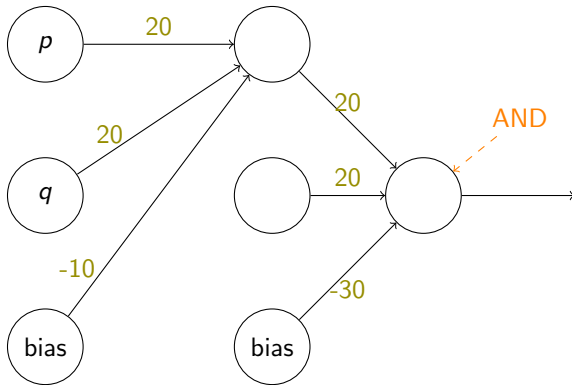
Computing 'xor'




Computing 'xor'



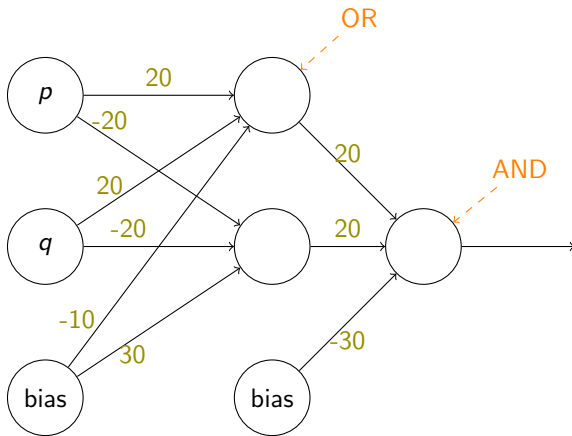
Computing 'xor'



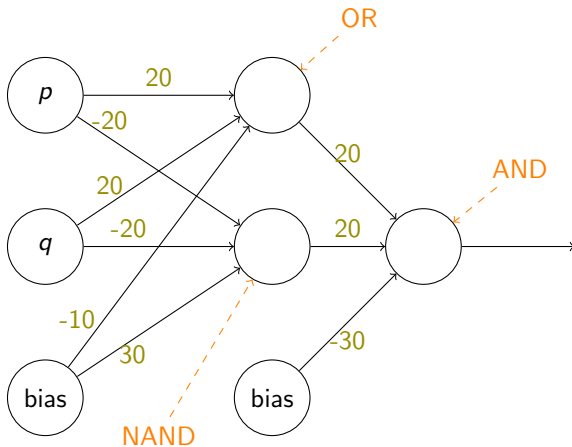
Computing 'xor'




Computing 'xor'



Computing 'xor'

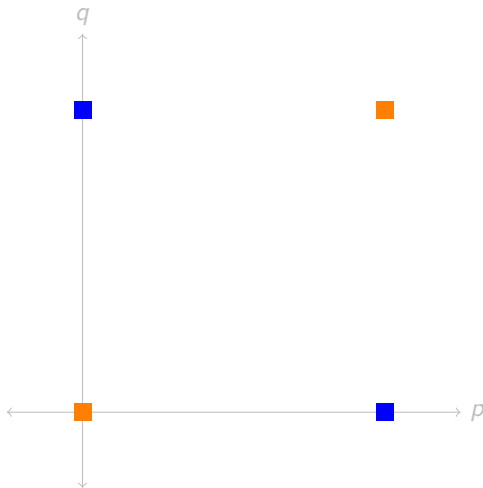


Computing 'xor'

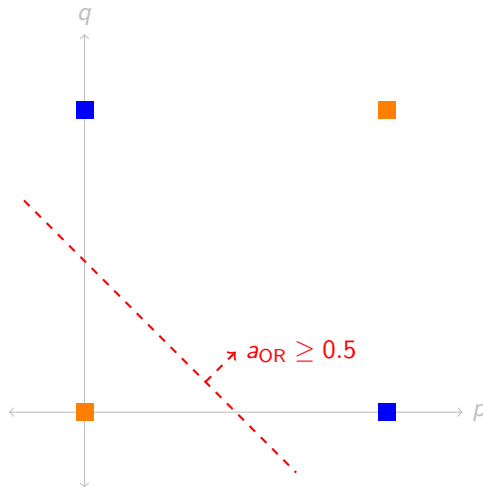


Exercise: show that the hidden units behave as labeled.


Computing 'xor'



Computing 'xor'



Computing 'xor'



Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\vec{a} = f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right)$$

$$= f(xW + b)$$

Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\vec{a} = f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right) \\ = f(xW + b)$$

- x_j^i : j th feature of input i
- w_k^l : weight from neuron k to neuron l in next layer
- b_m : bias to neuron m in next layer

Computing Many Examples

It's often very useful to compute over a 'batch' of inputs at once. The linear combination then turns into a *matrix multiplication*:

$$\vec{a} = f \left(\begin{bmatrix} x_0^0 & x_1^0 & \cdots & x_n^0 \\ x_0^1 & x_1^1 & \cdots & x_n^1 \\ \vdots & \vdots & \ddots & \vdots \\ x_0^m & x_1^m & \cdots & x_n^m \end{bmatrix} \begin{bmatrix} w_0^0 & w_0^1 & \cdots & w_0^l \\ w_1^0 & w_1^1 & \cdots & w_1^l \\ \vdots & \vdots & \ddots & \vdots \\ w_n^0 & w_n^1 & \cdots & w_n^l \end{bmatrix} + \begin{bmatrix} b_0 & b_1 & \cdots & b_l \\ b_0 & b_1 & \cdots & b_l \\ \vdots & \vdots & \ddots & \vdots \\ b_0 & b_1 & \cdots & b_l \end{bmatrix} \right)$$

$$= f(xW + b)$$

- x_j^i : j th feature of input i
- w_k^l : weight from neuron k to neuron l in next layer
- b_m : bias to neuron m in next layer

Exercises:

- write down W^1 and W^2 for the xor network.
- re-write the above as $f(xW)$ by adding a column of 1s to x and a new row to W .

Hidden Representations

Key idea: hidden layers of a neural network can encode high-level/abstract features of the input.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.
- Update the weights and biases to move the network's outputs closer to the true outputs.

(Supervised) Learning

Where do the weights and biases come from? They can be *learned* from data to approximate a function.

Supervised learning (will talk about others later):

- Initialize the network randomly.
- Give the network a bunch of inputs.
- Compare its outputs *to the true outputs*.
- Update the weights and biases to move the network's outputs closer to the true outputs.

The last step is done via *gradient descent* (and refinements thereof).

Gradient Descent: Example

Task: predict a true value $y = 2$.

Gradient Descent: Example

Task: predict a true value $y = 2$.

“Model”: one parameter θ , outputs $\hat{y} = \theta$.

Gradient Descent: Example


Task: predict a true value $y = 2$.

“Model”: one parameter θ , outputs $\hat{y} = \theta$.

Loss function:

$$\mathcal{L}(\theta, y) = (\hat{y}(\theta) - y)^2$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$


Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$

Gradient Descent: Example



$$\frac{\partial}{\partial \theta} \mathcal{L}(\theta, y) = 2(\theta - y)$$

$$\theta_{t+1} = \theta_t - \alpha \cdot \frac{\partial}{\partial \theta} \mathcal{L}(\theta, y)$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:

$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2) \cdots)W^L + b^L)$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:

$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2)\cdots)W^L + b^L)$$

All of the weights and biases form a long vector of parameters θ . So instead of a partial derivative, we take a *gradient*:

$$\nabla_{\theta} \mathcal{L}(\hat{y}(\theta), y) = \left\langle \frac{\partial}{\partial \theta_1} \mathcal{L}, \dots, \frac{\partial}{\partial \theta_N} \mathcal{L} \right\rangle$$

Gradient Descent for NNs

A neural network computes a complex function of its input. For an L -layer feed-forward network:

$$\hat{y}(x) = f_L(f_{L-1}(\cdots f_2(f_1(xW^1 + b^1)W^2 + b^2) \cdots)W^L + b^L)$$

All of the weights and biases form a long vector of parameters θ . So instead of a partial derivative, we take a *gradient*:

$$\nabla_{\theta} \mathcal{L}(\hat{y}(\theta), y) = \left\langle \frac{\partial}{\partial \theta_1} \mathcal{L}, \dots, \frac{\partial}{\partial \theta_N} \mathcal{L} \right\rangle$$

The (negative) gradient tells us *which direction in 'parameter space'* to walk in order to make the loss (\mathcal{L}) smaller, i.e. to make the network's output closer to the true output.

Learned Representations

Key idea: a neural network can learn *which* high-level/abstract features of the input are useful in helping it solve its task. (Features are learned, instead of engineered by us.)

Anatomy of a DL Experiment

① Specify parameters

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!
 - ① Evaluate at regular intervals
 - ② Measure your variables of interest
 - ③ Monitor train/test loss
 - ④ Early stopping [later in tutorial]
- ⑤ Analyze
 - quantitative
 - qualitative behavior

Anatomy of a DL Experiment

- ① Specify parameters
- ② Build data input/generation pipeline
 - Train/test split [dev as well; more later]
- ③ Build model
- ④ Train the model!
 - ① Evaluate at regular intervals
 - ② Measure your variables of interest
 - ③ Monitor train/test loss
 - ④ Early stopping [later in tutorial]
- ⑤ Analyze
 - quantitative
 - qualitative behavior

NOTE: keep detailed records about what you're doing!

For guidance on keeping records and writing up results, see [these lecture notes](#) from Sam Bowman.

To the code!

<https://github.com/shanest/nn-tutorial/blob/master/tutorial.ipynb>

Balancing the Data

For classification tasks, it's very important to *balance the training data*, so that it contains (roughly) equal numbers of examples for each class. Otherwise, a network can very quickly get stuck in a local minimum, where it uniformly guesses the most-frequent class.

Two methods:

- ① Downsampling: randomly sample as many examples from each class as in your least-frequent class
- ② Upsampling: repeat examples from your less-frequent classes until you have as many as the most-frequent

Early Stopping

How do you know how many epochs to train for? One common method:
A LOT, but put in a condition for *early stopping*.

Over-fitting: detection

Neural networks, especially very large, deep ones, run the risk of *over-fitting* their training data. (Sometimes, this is referred to as “memorizing” the training data.)

Over-fitting: avoidance (regularization)

Two very popular and successful *regularization* techniques to combat over-fitting (and generally make your life better):

- Dropout: randomly ‘turn off’ (set to 0) a certain percentage of input nodes to this layer
- Batch normalization: scale the inputs to the layer so that they’re roughly averaged around 0 and normally distributed.
(I’ve found BN *very* useful in my own research.)

Note: `.eval()` and `.train()` on a PyTorch `nn.Module` turn these sorts of methods off/on based on whether you are in training or ‘inference’ / prediction mode.

Hyper-parameter Tuning

How do you decide how to set the parameters of your experiment? There are so many knobs to turn! (Listed roughly in order of importance of turning them.)

- ① Network architecture: depth (number of layers) and width (size of layers)
[or even different network types altogether]
- ② Activation functions
- ③ Optimizers
 - Learning rates + other parameters here
- ④ ...

Hyper-parameter Tuning (cont.)

Main idea: try a whole bunch of settings! Choose the setting that performs best. Some notes:

- This requires a third set, a *development (dev)* set, in addition to training/testing.
You choose the best hyper-parameters based on best performance (however measured) on the *dev* set.
Then you evaluate that model on the *test* set.
- *Random* search in parameter space appears to be better than 'grid' search:
- This can be parallelized and semi-automated.

References I